

schulzdLab1

December 6, 2020

1 Lab 1: Data Cleaning

David Schulz

1.1 Introduction

In this lab, we are going to inspect and clean a data set of real estate transactions from California. We will determine when variables should be represented categorically, clarify what the variables represent, and filter variables down to make them more useful to our analysis.

1.2 1. Loading the Data and Initial Assessment

```
[1]: import pandas as pd

data = pd.read_csv("Sacramentorealestatetransactions.csv")

print(data.head())
print()
print(data.info())
```

	street	city	zip	state	beds	baths	sq__ft	\
0	3526 HIGH ST	SACRAMENTO	95838	CA	2	1	836	
1	51 OMAHA CT	SACRAMENTO	95823	CA	3	1	1167	
2	2796 BRANCH ST	SACRAMENTO	95815	CA	2	1	796	
3	2805 JANETTE WAY	SACRAMENTO	95815	CA	2	1	852	
4	6001 MCMAHON DR	SACRAMENTO	95824	CA	2	1	797	

	type	sale_date	price	latitude	longitude
0	Residential	Wed May 21 00:00:00 EDT 2008	59222	38.631913	-121.434879
1	Residential	Wed May 21 00:00:00 EDT 2008	68212	38.478902	-121.431028
2	Residential	Wed May 21 00:00:00 EDT 2008	68880	38.618305	-121.443839
3	Residential	Wed May 21 00:00:00 EDT 2008	69307	38.616835	-121.439146
4	Residential	Wed May 21 00:00:00 EDT 2008	81900	38.519470	-121.435768

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 985 entries, 0 to 984
```

```
Data columns (total 12 columns):
street      985 non-null object
city        985 non-null object
zip         985 non-null int64
state       985 non-null object
beds        985 non-null int64
baths       985 non-null int64
sq__ft      985 non-null int64
type        985 non-null object
sale_date   985 non-null object
price       985 non-null int64
latitude    985 non-null float64
longitude    985 non-null float64
dtypes: float64(2), int64(5), object(5)
memory usage: 92.4+ KB
None
```

Street is an object, city is an object, zip is an int64, state is an object, beds is an int64, baths is an int64, sq__ft is an int64, type is an object, sale_date is an object, price is an int64, latitude is a float64, and longitude is a float64. None of the columns have null values because they're all supposed to be non-null.

1.3 2. Representing Categorical Variables

```
[2]: print(data['street'].nunique())
      print(data['zip'].nunique())
      print(data['beds'].nunique())
```

981

68

8

Streets: 981

Zip codes: 68

Beds: 8

It's probably more appropriate to represent them as categorical variables because

```
[3]: data['city'] = data['city'].astype('category')
      data['state'] = data['state'].astype('category')
      data['zip'] = data['zip'].astype('category')
      data['beds'] = data['beds'].astype('category')
      data['baths'] = data['baths'].astype('category')
      data['type'] = data['type'].astype('category')

      print(data.info())
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 985 entries, 0 to 984
Data columns (total 12 columns):
street      985 non-null object
city        985 non-null category
zip         985 non-null category
state       985 non-null category
beds        985 non-null category
baths       985 non-null category
sq__ft      985 non-null int64
type        985 non-null category
sale_date   985 non-null object
price       985 non-null int64
latitude    985 non-null float64
longitude    985 non-null float64
dtypes: category(6), float64(2), int64(2), object(2)
memory usage: 57.5+ KB
None

```

1.4 3. Cleaning Continuous Variables


```

[4]: import matplotlib.pyplot as plt

plt.title('# of Times sq__ft in Each Group')
plt.xlabel('sq__ft')
plt.ylabel('# of Occurrences')
plt.hist(data['sq__ft'], bins=20)

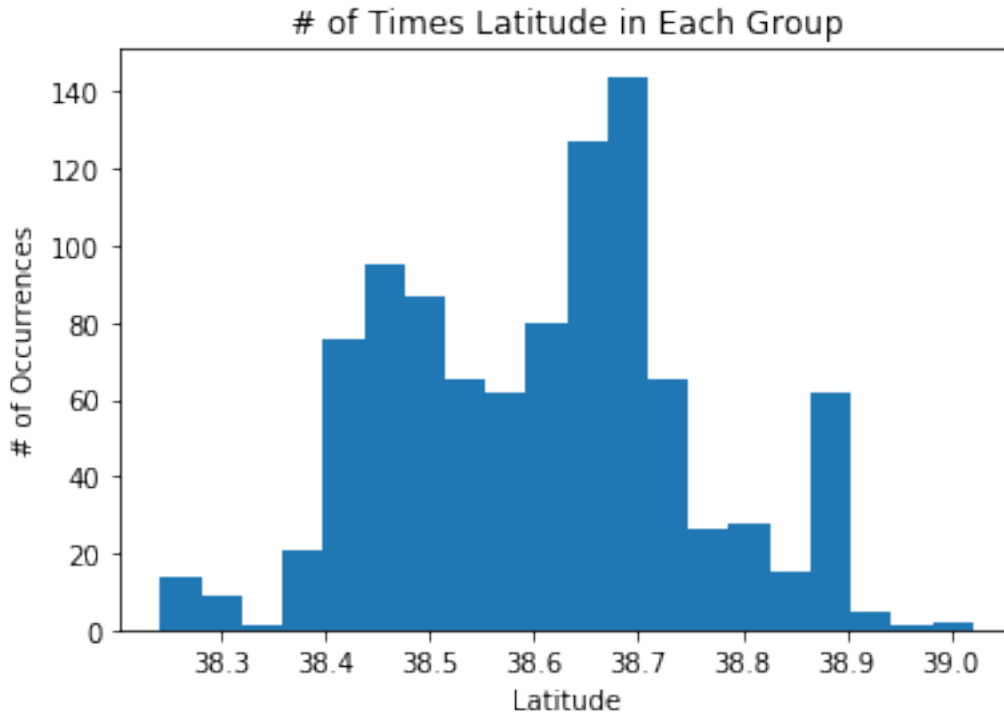
[4]: (array([171.,  2., 46., 175., 210., 145., 82., 56., 36., 19., 13.,
          11.,  9.,  6.,  2.,  1.,  0.,  0.,  0.,  1.]),
      array([ 0. , 291.1, 582.2, 873.3, 1164.4, 1455.5, 1746.6, 2037.7,
          2328.8, 2619.9, 2911. , 3202.1, 3493.2, 3784.3, 4075.4, 4366.5,
          4657.6, 4948.7, 5239.8, 5530.9, 5822. ]),
      <a list of 20 Patch objects>)

```



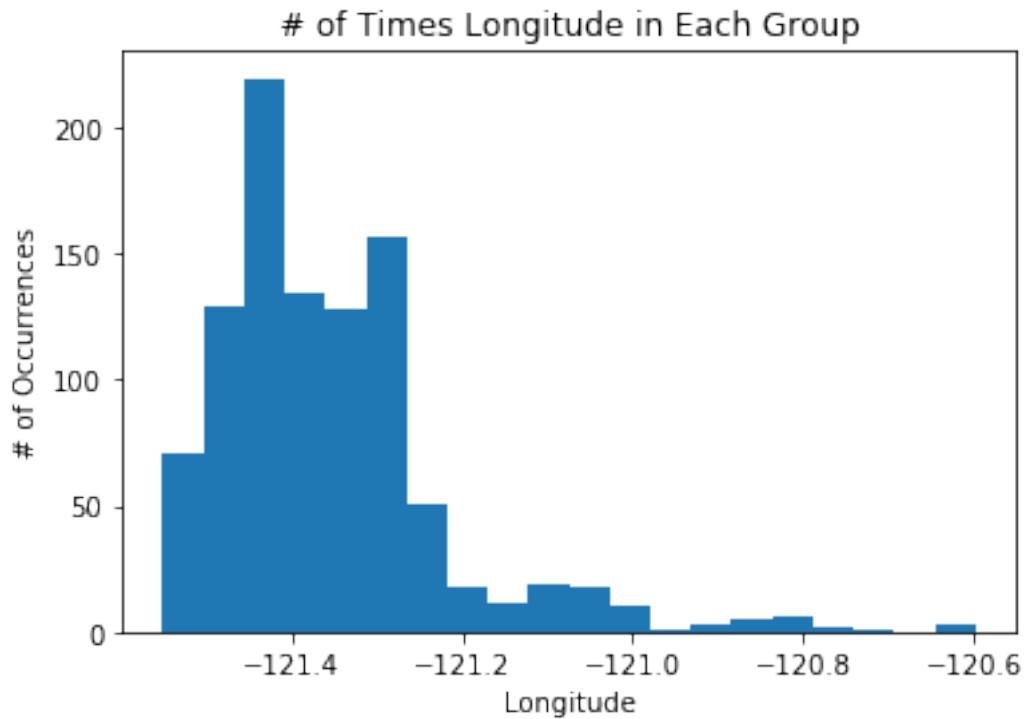
```
[5]: plt.title('# of Times Latitude in Each Group')
plt.xlabel('Latitude')
plt.ylabel('# of Occurrences')
plt.hist(data['latitude'], bins=20)
```

```
[5]: (array([ 14.,   9.,   1.,  21.,  76.,  95.,  87.,  65.,  62.,  80., 127.,
        144.,  65.,  26.,  28.,  15.,  62.,   5.,   1.,   2.]),
      array([38.241514, 38.2804787, 38.3194434, 38.3584081, 38.3973728,
        38.4363375, 38.4753022, 38.5142669, 38.5532316, 38.5921963,
        38.631161, 38.6701257, 38.7090904, 38.7480551, 38.7870198,
        38.8259845, 38.8649492, 38.9039139, 38.9428786, 38.9818433,
        39.020808 ]),
      <a list of 20 Patch objects>)
```



```
[6]: plt.title('# of Times Longitude in Each Group')
plt.xlabel('Longitude')
plt.ylabel('# of Occurrences')
plt.hist(data['longitude'], bins=20)
```

```
[6]: (array([ 71., 129., 219., 134., 128., 156.,  51.,  18.,  11.,  19.,  18.,
          10.,   1.,   3.,   5.,   6.,   2.,   1.,   0.,   3.]),
      array([-121.551704, -121.50399875, -121.4562935, -121.40858825,
          -121.360883, -121.31317775, -121.2654725, -121.21776725,
          -121.170062, -121.12235675, -121.0746515, -121.02694625,
          -120.979241, -120.93153575, -120.8838305, -120.83612525,
          -120.78842, -120.74071475, -120.6930095, -120.64530425,
          -120.597599 ]),
      <a list of 20 Patch objects>)
```



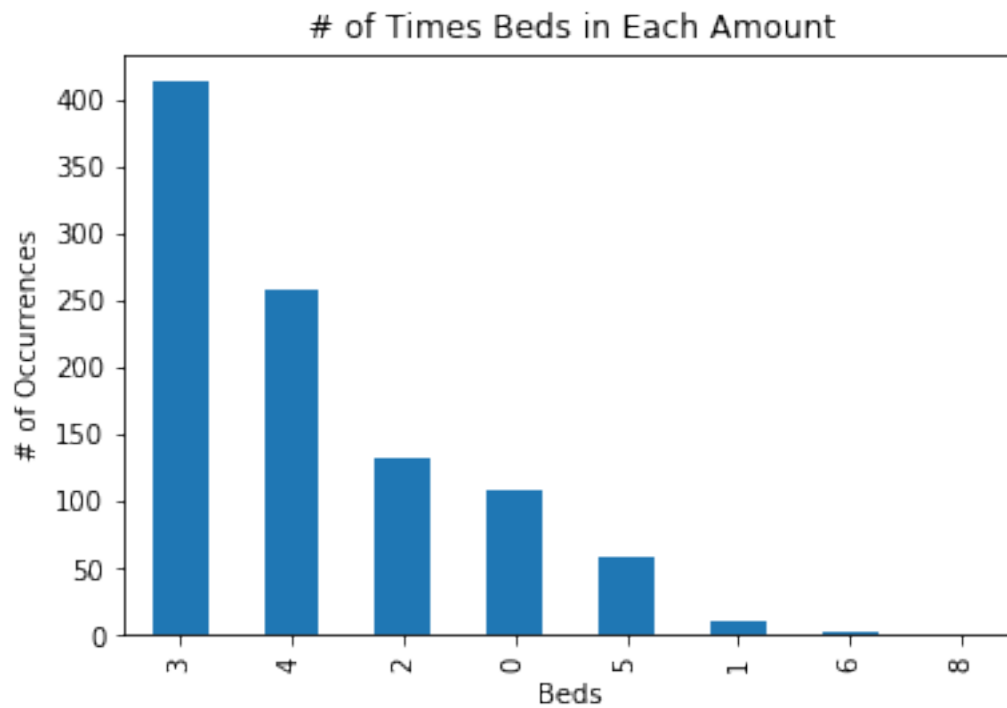
A histogram with a sufficient number of bins to show the number of entries in a small group accurately enough is appropriate.

The fact that almost 175 of the entries have a square footage of zero is very noticeable. Since it's impossible to have a square footage of zero, I'm pretty sure they are artifacts.

1.5 4. Cleaning Categorical Variables

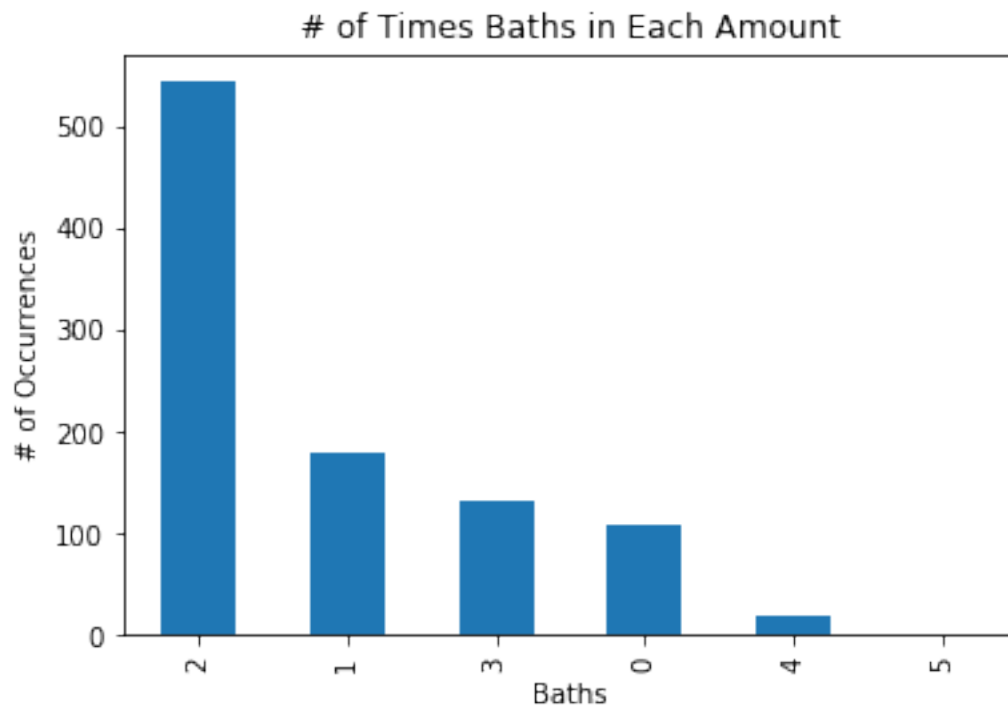
```
[7]: plt.title('# of Times Beds in Each Amount')
plt.xlabel('Beds')
plt.ylabel('# of Occurrences')
data['beds'].value_counts().plot(kind='bar')
```

```
[7]: <matplotlib.axes._subplots.AxesSubplot at 0x7f01808aaad0>
```




```
[8]: plt.title('# of Times Baths in Each Amount')  
plt.xlabel('Baths')  
plt.ylabel('# of Occurrences')  
data['baths'].value_counts().plot(kind='bar')
```

```
[8]: <matplotlib.axes._subplots.AxesSubplot at 0x7f01807db510>
```




```
[9]: plt.title('# of Each Type')  
plt.xlabel('Type')  
plt.ylabel('# of Occurrences')  
data['type'].value_counts().plot(kind='bar')
```

```
[9]: <matplotlib.axes._subplots.AxesSubplot at 0x7f0180748e90>
```

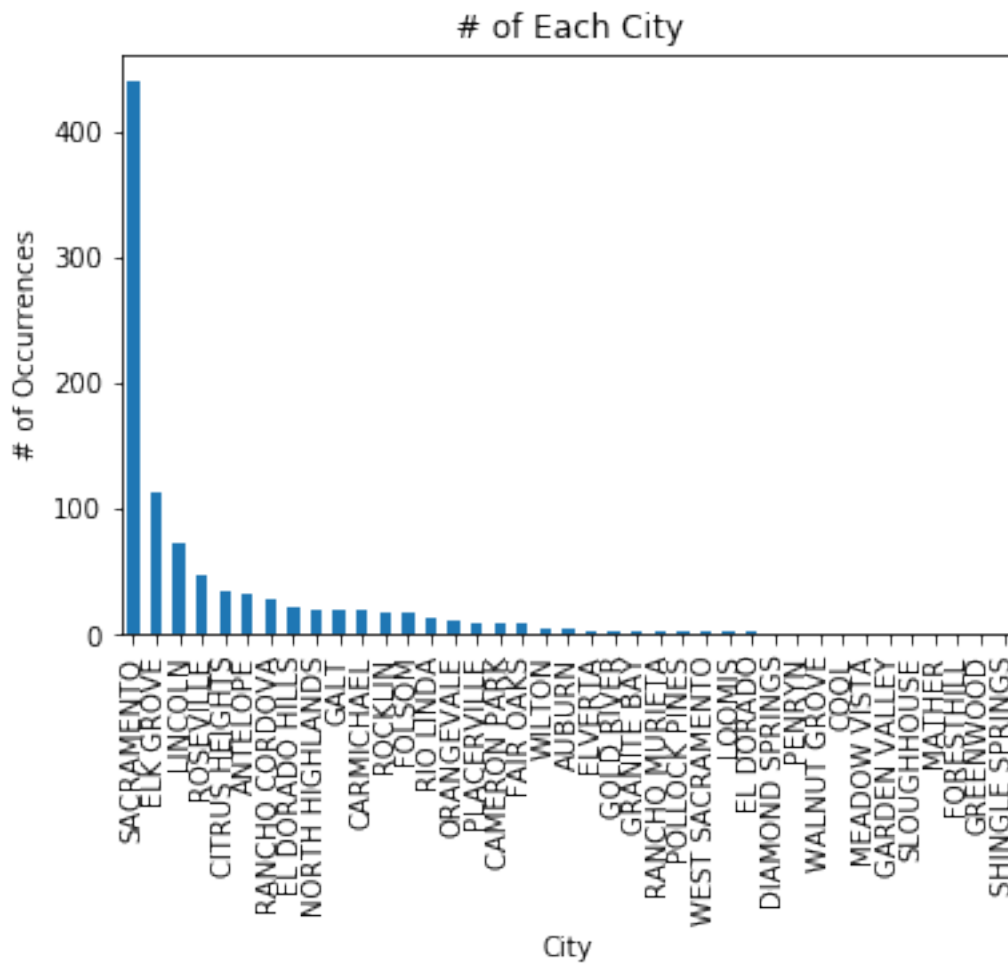
```
[10]: plt.title('# of Each State')  
plt.xlabel('State')  
plt.ylabel('# of Occurrences')  
data['state'].value_counts().plot(kind='bar')
```

```
[10]: <matplotlib.axes._subplots.AxesSubplot at 0x7f0180726b90>
```




```
[11]: plt.title('# of Each City')
      plt.xlabel('City')
      plt.ylabel('# of Occurrences')
      data['city'].value_counts().plot(kind='bar')
```

```
[11]: <matplotlib.axes._subplots.AxesSubplot at 0x7f01807322d0>
```



```
[12]: plt.title('# of Each Zip Code')
plt.xlabel('Zip Code')
plt.ylabel('# of Occurrences')
data['zip'].value_counts().plot(kind='bar')
```

```
[12]: <matplotlib.axes._subplots.AxesSubplot at 0x7f0180689950>
```



There are properties that have 0 bedrooms and/or 0 bathrooms. They could be empty lots that have yet to be built on.

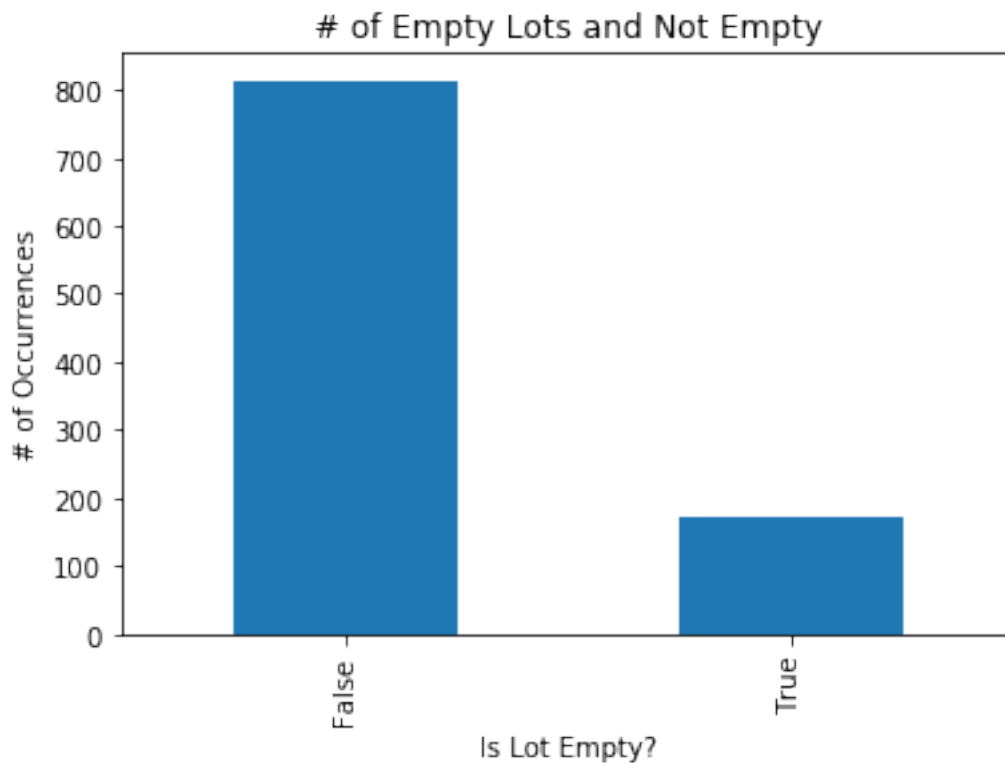
1.6 5. Engineering New Variables - Part I

```
[13]: import numpy as np

data['empty_lot'] = np.where(data['sq_ft'] == 0, True, False)

plt.title('# of Empty Lots and Not Empty')
plt.xlabel('Is Lot Empty?')
plt.ylabel('# of Occurrences')
data['empty_lot'].value_counts().plot(kind='bar')
```

```
[13]: <matplotlib.axes._subplots.AxesSubplot at 0x7f01803ed990>
```



1.7 6. Engineering New Variables - Part II

```
[14]: print(data['street'].nunique())
```

981

It's not useful for analysis or ML in its current form.

```
[15]: data.head(n=20)
```

```
[15]:
```

	street	city	zip	state	beds	baths	\
0	3526 HIGH ST	SACRAMENTO	95838	CA	2	1	
1	51 OMAHA CT	SACRAMENTO	95823	CA	3	1	
2	2796 BRANCH ST	SACRAMENTO	95815	CA	2	1	
3	2805 JANETTE WAY	SACRAMENTO	95815	CA	2	1	
4	6001 MCMAHON DR	SACRAMENTO	95824	CA	2	1	
5	5828 PEPPERMILL CT	SACRAMENTO	95841	CA	3	1	
6	6048 OGDEN NASH WAY	SACRAMENTO	95842	CA	3	2	
7	2561 19TH AVE	SACRAMENTO	95820	CA	3	1	
8	11150 TRINITY RIVER DR Unit 114	RANCHO CORDOVA	95670	CA	2	2	
9	7325 10TH ST	RIO LINDA	95673	CA	3	2	
10	645 MORRISON AVE	SACRAMENTO	95838	CA	3	2	

11	4085 FAWN CIR	SACRAMENTO	95823	CA	3	2
12	2930 LA ROSA RD	SACRAMENTO	95815	CA	1	1
13	2113 KIRK WAY	SACRAMENTO	95822	CA	3	1
14	4533 LOCH HAVEN WAY	SACRAMENTO	95842	CA	2	2
15	7340 HAMDEN PL	SACRAMENTO	95842	CA	2	2
16	6715 6TH ST	RIO LINDA	95673	CA	2	1
17	6236 LONGFORD DR Unit 1	CITRUS HEIGHTS	95621	CA	2	1
18	250 PERALTA AVE	SACRAMENTO	95833	CA	2	1
19	113 LEEWILL AVE	RIO LINDA	95673	CA	3	2

	sq__ft	type	sale_date	price	latitude	\
0	836	Residential	Wed May 21 00:00:00 EDT 2008	59222	38.631913	
1	1167	Residential	Wed May 21 00:00:00 EDT 2008	68212	38.478902	
2	796	Residential	Wed May 21 00:00:00 EDT 2008	68880	38.618305	
3	852	Residential	Wed May 21 00:00:00 EDT 2008	69307	38.616835	
4	797	Residential	Wed May 21 00:00:00 EDT 2008	81900	38.519470	
5	1122	Condo	Wed May 21 00:00:00 EDT 2008	89921	38.662595	
6	1104	Residential	Wed May 21 00:00:00 EDT 2008	90895	38.681659	
7	1177	Residential	Wed May 21 00:00:00 EDT 2008	91002	38.535092	
8	941	Condo	Wed May 21 00:00:00 EDT 2008	94905	38.621188	
9	1146	Residential	Wed May 21 00:00:00 EDT 2008	98937	38.700909	
10	909	Residential	Wed May 21 00:00:00 EDT 2008	100309	38.637663	
11	1289	Residential	Wed May 21 00:00:00 EDT 2008	106250	38.470746	
12	871	Residential	Wed May 21 00:00:00 EDT 2008	106852	38.618698	
13	1020	Residential	Wed May 21 00:00:00 EDT 2008	107502	38.482215	
14	1022	Residential	Wed May 21 00:00:00 EDT 2008	108750	38.672914	
15	1134	Condo	Wed May 21 00:00:00 EDT 2008	110700	38.700051	
16	844	Residential	Wed May 21 00:00:00 EDT 2008	113263	38.689591	
17	795	Condo	Wed May 21 00:00:00 EDT 2008	116250	38.679776	
18	588	Residential	Wed May 21 00:00:00 EDT 2008	120000	38.612099	
19	1356	Residential	Wed May 21 00:00:00 EDT 2008	121630	38.689999	

	longitude	empty_lot
0	-121.434879	False
1	-121.431028	False
2	-121.443839	False
3	-121.439146	False
4	-121.435768	False
5	-121.327813	False
6	-121.351705	False
7	-121.481367	False
8	-121.270555	False
9	-121.442979	False
10	-121.451520	False
11	-121.458918	False
12	-121.435833	False
13	-121.492603	False

```
14 -121.359340      False
15 -121.351278      False
16 -121.452239      False
17 -121.314089      False
18 -121.469095      False
19 -121.463220      False
```


The street type is almost always the last word of the string, unless it's followed by "Unit".

```
[16]: def get_street_type(address) -> str:
      words = address.split()
      if "Unit" in address:
          return words[-3]
      elif "HIGHWAY" in address:
          return words[-2]
      elif "VIA" in address:
          return "WAY"
      elif "AVENIDA" in address:
          return "AVE"
      elif "MADERA" in address:
          return "VIEW"
      else:
          return words[-1]
```

```
[17]: data['street_type'] = data['street'].map(get_street_type)

      plt.title('# of Each Street Type')
      plt.xlabel('Street Type')
      plt.ylabel('# of Occurrences')
      data['street_type'].value_counts().plot(kind='bar')
```

```
[17]: <matplotlib.axes._subplots.AxesSubplot at 0x7f0180408050>
```



1.8 7. Identifying Potential Dependent Variables

- Variables that are continuous, such as floats, are appropriate for regression.
- Variables that are not continuous and have specific values from a given list of possible choices, such as strings or integers, are appropriate for classification.
- The average speed, in MPH, of a vehicle, which would be a float, would be a good dependent variable for a regression problem.
- The make of a vehicle, which would be a string chosen from a small list of possible strings, would be a good dependent variable for a classification problem.

1.9 8. Save the Cleaned Data Set

```
[18]: data.to_csv("cleaned_Sacramentorealestatetransactions.csv", index=False)
```